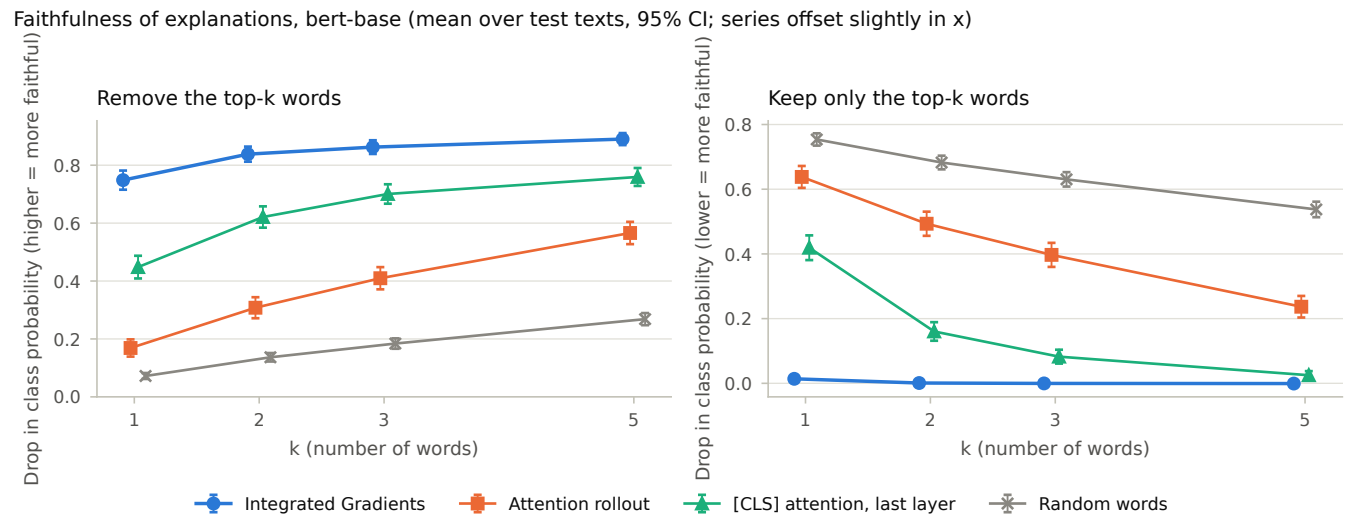
Faithfulness of explanations, bert-base (mean over test texts, 95% CI; series offset slightly in x)
Remove the top-k words
Keep only the top-k words
Drop in class probability (higher = more faithful)
Drop in class probability (lower = more faithful)
0.0
0.2
0.4
0.6
0.8
1
2
3
5
k (number of words)
Integrated Gradients
Attention rollout
[CLS] attention, last layer
Random words